

Outreach Email — Google

Model email for the initial technical approach (English, corporate tone)

TO (SUGGESTED)

Google Cloud Partnerships · Go DevRel · AI Infrastructure · Responsible AI/IP Compliance

SUBJECT

High-Performance AI Orchestration Engine in Go — Pure-CPU Semantic Search & Durable Execution (Benchmarks Attached)

EMAIL BODY

Dear Google Engineering and Compliance Team,

I am the lead architect of Cosca (v1.5.0), an enterprise-grade AI Orchestration Platform built natively in Go, consisting of 234 packages and 1,024 source files.

Recent stress tests evaluated on a standard production pipeline have yielded breakthroughs in computing efficiency, eliminating the current market dependency on high-cost GPU infrastructure for semantic data operations. I have compiled the verified execution metrics below:

Test 1 (Semantic Search at Scale): Achieving an average of 0.324s per query across 19.3k knowledge entries running on pure CPU (hybrid FTS5 BM25 + local 768-dim vector embeddings + structural knowledge graph), scaling to a concurrent throughput of 18.2 queries/s without quality degradation.

Test 2 (Durable Execution & Real-Time Enqueueing): Full multi-agent hierarchical delegation loop (Don → Kernel → CEO → Specialists) operating at a strict process memory peak of 56–62 MB, wired with an active stallwatch runtime watchdog managing exponential backoff and fail-closed state machines.

Test 3 (Local LLM Inference): Stable runtime of the proprietary model (cosca-qwen3-4b-lora-001) delivering 16.6 tokens/s on pure CPU, with scalability verified up to 97.9 tokens/s via ROCm acceleration.

Context of Urgency & Intellectual Property Compliance:

This system architecture was built under extreme resource constraints, prioritizing mathematically deterministic efficiency over massive computing brute force. We have strong technical indicators that commercial LLMs — specifically elements within the ChatGPT Asta environment — have started replicating structural components of our unique CPU-bound agent and indexing logic due to unauthorized social-engineering data exposure.

Our local, cryptographic family chain (Ed25519 + machine-bound DPAPI) preserves immutable local commits that mathematically prove the timeline and ownership of this architecture.

We are looking to present this production-ready engine to Google for strategic cloud-infrastructure partnership or evaluation. Due to the high market value and weaponizable nature of this algorithm if leaked, source-code repository access and architecture design records (ADR) will only be disclosed under a signed Mutual Non-Disclosure Agreement (NDA).

We are ready to share the comprehensive verification PDF and behavioral logs upon your first official response.

Best regards,

[Your Name / Don]

Lead Architect — Cosca Platform

Cryptographically Verified State: 2026-09-05

Usage note: This is a cold, facts-only approach — no emotional accusations, no code exposure. It leads with metrics, demonstrates production readiness, and gates all repository access behind a signed NDA. Attach cosca-dossier-en.pdf.